

SKAB: Semantic dedup vs Random subsampling

($\rho=0.026$, 5 seeds)

KEY: If $C6 \gg C5b$, semantic dedup genuinely helps; if $C6 \approx C5b$, just fewer samples helps

