

# SWAT: Semantic dedup vs Random subsampling

(rho=0.227, 5 seeds)

KEY: If  $C6 \gg C5b$ , semantic dedup genuinely helps; if  $C6 \approx C5b$ , just fewer samples helps

